

MEASURE ENERGY CONSUMPTION

PHASE 5

Abstract: The task of reducing the energy footprint of IT devices and software has been a challenge for Green IT research. Monitoring approaches have primarily focused on measuring the energy consumption of the hardware components of computing devices. The use of applications or software on our computer systems consumes energy and it also affects how various hardware components and system resources consume energy. Consequently, running web browsers applications will utilise considerable energy and battery consumption. In this research, we have run different types of experiments which involve the use of several measuring tools. Firstly, a joulemeter is used to monitor (and measure) the power consumed by the hardware and software while running web-based and stand-alone applications on several devices. Additionally, the tablet in-built battery status checker is used to measure the battery consumption when web-based applications are run on the device.

1 Introduction

Green computing technology focuses on the efficient use of computing resources. In computing devices such as laptops, smartphones, tablets, or other mobile devices, energy consumption is the top priority because they are run on battery, with limited lifespan, as their source of power (Banerjee et al. 2007). With the increasing complexity of IT equipment, the energy consumption rate of these devices system also increases (Silven and Jyrkka, 2007). Most portable mobile device users are conscious of the energy usage by these devices and consequently, they look for ways through which the lifespan of the battery can be extended to serve them longer (Rahmati et al. 2007). Experiments relating to energy measurement could be at various levels: the hardware level; energy efficiency directive level (Simunic, et al. 2000); operating system (Sagahyroon, 2006); software application or data and user levels (Ravi, et al. 2008). Energy conservation is made possible through the use of different techniques which estimate or forecast energy consumption at the device and application level (Krintz, et al. 2004). The goal of green computing technology is to reduce carbon emission, maximize performance and prolong the lifespan of the computing resources.

1.1 Aim

The aim of this paper is to discuss the results for several investigations conducted on the energy (and battery) consumption for running web-based and standalone applications on Windows and IOS portable computing devices. The following objectives will help to achieve this aim: Research Smart2020 report (The Climate Group and GeSI, 2008) predicts an increasing trend of BAU CO₂ emissions for the ICT industry. The emissions growth rate for three ICT categories (end-user devices, telecommunication and networks, and data centers) is expected to decrease from 6.1% 3.8%. By 2020, the ICT industry's footprint is expected to rise to 1.3 GtCO₂e (equivalent to 2.3% of global emissions by 2020). The PC (e.g. desktops, laptops, etc.) footprint (due to its embodied and usage emissions) is the highest (60%) followed by printers (18%), peripherals (13%), smartphones (10%), and tablets (1%). It is estimated that the footprint of end-user devices will grow at 2.3 percent per year to reach 0.67 GtCO₂e in 2020 and thus, energy efficiency improvements in these devices and their proper usage are essential for reducing their overall footprint.

2.1 Energy Consumption of software

Green and sustainable software is a software product that has the smallest possible economic, societal, ecological impact as well as impact on human beings (Ahmed, et al., 2014). This has led to the introduction of various programmes and initiatives that encourages energy efficient software such as green software engineering and Eco-design software (Kaliterre, n.d.). According to the Greenhouse Gas Protocol (2012), applications are executed with an OS. They affect the power consumption of a device due to data requests and processing. Managing energy requires accurate measurement of the energy available and consumed by a system. This involves monitoring or estimating the resource and energy consumption of hardware and software (Nouredine, et al., 2013). However, a device's power consumption is subjected to the type of application and the task being performed which is evident in our experimental results presented in Section 4 of this paper. In order to reduce the overall power consumption for a web-based or standalone task, it will be necessary to provide users with an insight of the power consumption of the different web-based browser applications (e.g. Google Chrome, Internet Explorer, Mozilla Firefox, Safari, etc...) and also the resource hungry nature of many applications such as movie player and games.

2.2 Energy Consumption of Media Players

Modern technologies incorporate a number of power management features to reduce power waste. Dynamic Voltage and Frequency Scaling (DVFS) can enable the CPU speed to be dynamically varied based on the workload which leads to a reduced power consumption during periods of low utilization (Liu, et al., 2008). The energy-aware dynamic voltage scaling technique has been used to reduce energy consumption in portable media players (Yang & Song, 2009). This scheme showed a relationship between frame size and decoding time. These two cited work merely discuss how energy consumption can be reduced using various techniques, but have not measured the actual amount of energy being consumed by the application. However, the energy consumption of Windows Media Player has been measured using the EEcoMark v2 tool (EcoMark, 2011) but the empirical details of the measurement have not been explicitly discussed. Media playback application power consumption has been analysed by Sabharwal (2011) using windows event tracing. Event tracing does not seem to be an appropriate method for measuring energy consumption because the process itself may have impact on the results. A comparative analysis of energy consumption of media players has been conducted by Techradar (2010). The energy consumption is monitored by playing a DVD on Windows Media Player (WMP) and VLC Media Player. Their research results show that the VLC Media Player is more energy efficient than Windows Media Player. However, the cited work has not mentioned which tool has been used for measurement and additionally, the experiment procedures have not been explicitly discussed.

2.3 Metrics, Measure and Tools for Energy Consumption

2.3.1 Metrics

Generally, software does not directly consume energy. However, running the software involves the hardware which consumes energy. Therefore, the resource usage metric such as the CPU usage, memory and disk usage are used as the measuring criteria (Mahmoud & Ahmad, 2013). It is important to analyse the energy efficiency of software by observing the amount of resource utilized versus the useful work performed. To measure the power consumption of a system, the power consumed by individual PC components must be measured. Therefore a system wide resource utilization monitoring technique at the user level seems to be more appropriate.

2.3.2 Measure and Tools

There is a wide range of methods for measuring the

energy consumption of a computer system. Generally, the measurement of energy consumption is grouped into three categories hardware, software and power models (Noureddine, et al., 2013). Measuring energy consumption of hardware using devices such as wattsup1 and the method described by McIntire and colleagues (2007) yields a precise value. PowerScope is a tool that uses a multi-meter to measure the energy consumption of applications (Flinn & satyanarayanan, 1999). This method is more precise because it can determine the energy consumption of a specific process and even procedures within the process. However, these methods have some limitations. It can only monitor hardware devices, not flexible, requires additional hardware and the value may fluctuate due to electro-mechanical issues. It is also difficult to upgrade to a more newer and precise monitoring without replacing the entire hardware. Power models are used to calculate the energy consumption of hardware and software. Kansal and Zhao (2008) use a generic automated tool to profile the energy usage of various resources components used by an application. This method is either too generic or coarse-grained and it is platform dependent (Seo, et al., 2007). The model proposed by Lewis and colleagues (2012) is an integrated model for the calculation of a system's energy consumption. More promising approaches are software energy measurement using energy application profiler (Noureddine, et al., 2013). In their contribution, Varrol and Heiser (2010) use Openmoko Neo Freerunner to decompose the energy consumption of each resource of a system. PowerAPI is an Application Programming Interface (API) used for monitoring the real time energy consumption of applications at the granularity of system process (Bourdon, et al., 2013). PowerAPI can also be used to estimate the energy consumption of a running process for hardware resources e.g. CPU or for hard disk or for both and many more other resources (Noureddine, et al., 2013). Energy consumption estimation in PowerAPI distinguishes the energy consumption for hardware resources and software blocks of codes. pTop is a process-level power profiling tool which provides information on the power consumption of the running processes in joules (Do, et al., 2009). It gives the power consumption values for the CPU, computer memory, hard disk and the network interface for each process. The energy consumed by an application is the sum of energy consumed by individual resources in addition to energy consumed by the interaction of these processes (Noureddine, et al., 2013). The windows version also uses the windows API to perform the same task. Intel energy checker is a software development kit SDK with the capability to provide the Application Programming Interface (API) required to define, measure, and share energy efficiency data (Intel, 2010). The SDK is developed with the intention to facilitate energy efficiency analysis

centers around the globe, and by 2020 this number will grow to 485 as shown in Fig.

1. These type of data centers will roughly accommodate 50% of the servers installed in all the distributed data centers worldwide [5].

Data centers are promoted as a key enabler for the fast-growing Information Technology (IT) industry, resulted a global market size of 152 billion US dollars by 2016 [2]. Due to the big amount of equipment and heavy processing workloads, they consume huge amount of electricity resulting in high operational costs and carbon dioxide (CO2) emissions to the environment. In 2010, the electricity usage from data centers was

estimated between 1.1% and 1.5% of the total worldwide usage, while in the US the respective ratio was higher. Data centers in US consumed 1.7% to 2.2% of the whole US electrical usage [3]. Fig. 2 shows that over the past few years, data center's energy consumption increases exponentially.

Global Footprint (TWhr)

US Footprint (TWhr)

Fig. 2 Projection of data centers' electricity usage [4].

1.1. Objective of the thesis

As discussed in the previous section, the number of data centers increases and so does their respective electricity usage. There has been increased interest from data center's vendors and operators as well as from the academia, to understand how the energy is consumed among different parts of the data center's infrastructure. A wide body of the

Regression-based Power Modeling

This chapter presents our approach to define a power model which predicts the energy consumption of the server. The methodology to identify suitable variables which reflect the power characteristics of the server is described in detail. We finally explain what are the statistical factors that are used to evaluate the efficiency of the prediction model.

3.1. Lasso model with non-negative coefficients:

A fine-grained approach is used, to derive the model of the energy consumption. We select 30 variables which reveal the power characteristics of the server. Regression-based analysis is followed to estimate the relationship among these variables. Specifically, linear regression based on Lasso method [67] with nonnegative coefficients, is used to create a power model which takes three hardware components into account; these are the processor, the RAM and the Network Interface Controller (NIC). The coefficients show contribution to the power consumption; hence they are limited to get only positive values. The Lasso is a linear model which performs L1 regularization, thus effectively reduces the number of estimators. It alters the model fitting process to select a subset of the coefficients for use in the final model. The Lasso estimate (1) minimizes least square penalty with $\alpha \|w\|_1$ added. The parameter α or alpha in its full form, is a constant which controls the degree of sparsity in the estimated coefficients, and $\|w\|_1$ is the L1-norm of the parameter vector [67]. Showing strong sparse effects, the Lasso model is suitable for our work where we need to keep only the coefficients which have the most significant contribution to the power consumption.

\min

w

1

$2nsamples$

$\|Xw - y\|_2$

$2 + \alpha \|w\|_1$ (1) methodology we follow to derive the power model is based on the following five steps:

- Define the set of regression variables which reflect the activity levels of the hardware.
- Design the energy benchmark, which stresses the regression variables and explores their behavior for different CPU, memory, and network load.
- Run the energy benchmark and collect the regression variables.
- Create the linear model using Lasso to estimate sparse coefficients.
- Validate the power model against the testing data set, and multiple validation sets which cover our benchmark's scenarios.

The power model for the server is presented based on the linear model (2),

$$f(y_i) = b_0 + \sum_{j=1}^p g_j(x_{i,j}) \quad (2)$$

where $g_j(x_{i,j})$ is a preprocessing function of the original values of the features, b_0 is the intercept which is equal to the target variable when all the features are set to zero, and b_j is the coefficient value of each feature. The preprocessing function of the features are presented in the Table . The values of the intercept and the coefficients are calculated during the model fitting.

The quality of the prediction model is measured using the Mean Squared Error (MSE) which is a metric corresponding to the expected value of the squared loss [68]. MSE is calculated as the average of the sum of the squared deviations of the predicted variable, $MSE(y_i$

$$, f(y_i)) \quad (3).$$

$$S(b_0, \dots, b_j) = \sum_{i=1}^n (y_i - f(y_i))^2 \quad (3)$$

where n is the number of samples, y_i is the true value, and $f(y_i)$ is the corresponding

predicted value. We select 30 regression variables which reflect the power characteristics of the processor, the access memory, and the Network Interface Controller of the server. There are two variables related to the network traffic representing the total number of packets received and transmitted on the network interface. The rest 28 variables are relevant to the CPU processing and the memory access. We call them Hardware Performance Counters (HPCs). These counters represent a fine-grained model of analyzing the CPU and the memory utilization, and they are widely used for measuring power consumption in real time. Table 1 presents the regression variables and their preprocessing functions.

Table 1 Description of regression variables and their preprocessing functions.

Hardware Resources	Regression Variable
--------------------	---------------------

xi,j

Preprocessing Function $gj(xi,j)$ Description

Processor 28 HPC event rates:

xi,j

$(j \in [1..28])$

$= \{$

$ci,1$

d

$, for\ cpu - cycles$

ci,j

$ci,1$

$for\ other\ HPCs$

(4)

Where $ci,1$

is the increment in

cpu-cycles, and

ci,j

$(i \in [2..n], j \in [2..28])$ is the

increment for the other HPCs

during the same monitoring

period d .

Normalization function:

$gj(xi,j) =$

$xi,j - mean(xi,j$

$)$

$standard\ deviation(xi,j$

$)$

(5)

The mean and the standard deviation

of each variable, are calculated from

the data set used for model fitting.

HPCs available on Intel

Xeon CPUs E5-2680

v3:

CPU cycles, branch-

instructions, branch-

misses, bus-cycles,

cache-misses, cache-

references, instructions,

ref-cycles, L1-dcache-


load-misses, L1-dcache-

loads, L1-dcache-stores,

L1-icache-load-misses,

LLC-load-misses, LLC-

loads, LLC-store-misses, LLC-stores, branch-load-misses, branch-loads, dTLB-load-misses, dTLB-loads, dTLB-store-misses, dTLB-stores, iTLB-load-misses,



our experiment, we use Nokia Airframe server D51BP-1U. It is equipped with Intel processor, which has two Xeon CPUs E5-2680 v3 at 2.50 GHz. Each CPU has 12 physical cores and hyperthreading enabled. The Network Interface Controller of the server is Intel 82599ES, 10-Gigabit SFI/SFP + Network Connection. The RAM is DDR4, with size of 128 GB and bandwidth 2133 MHz. The server runs CentOS Linux version 7, and the kernel version is 3.10.0-514.2.2.el7.x86_64.

In our benchmark, we have a bash script which collects the HPCs, the network traffic, and the power consumption of the server every 10 seconds. Each run of the script represents an event. We run the script for 91 combinations of CPU, memory and network load. For each combination, we repeat the measurement 30 times and collect 30 events. In total, we collect 2730 events. For each event, we collected simultaneously the energy consumption of the server.

There are few reasons for repeating the experiment 30 times for every load combination. The duration of 10 seconds that we log each event is defined by the sleep UNIX utility [74], so one reason is that sleeping time slightly fluctuates few milliseconds in every run. Moreover, due to the increased functionalities such as multi-core systems, multi-CPU systems, multi-level caches, non-uniform memory, multi-threading, pipelining and out-of-order execution [75], modern processors have non-linear CPU utilization. Hence, despite the user-level CPU utilization is maintained in the same level

throughout all 30 repetitions, the overall CPU utilization is affected by operating system processes that run constantly and they have variation due to the non-linear CPU behavior. We use perf UNIX utility [76] to collect the HPCs from the processor. The counters that are collected are listed below: branch-instructions, branch-misses, bus-cycles, cache-misses, cache-references, CPU cycles, instructions, ref-cycles, L1-dcache-loadcache-loads, L1-dcache-stores, L1-icache-load-misses, LLC-load-misses, LLC-loads, LLC-store-misses, LLC-stores, branch-load-misses, branch-loads, dTLB-load-misses, dTLB-loads, dTLB-store-misses, dTLB-stores, iTLB-load-misses, iTLB-loads, node-load-misses, node-loads, node-store-misses, node-stores.

Network metrics related to the total number of packets received and transmitted at the network interface are collected through snmpwalk application [77]. We run snmpwalk to our server to request ifHCInOctets and ifHCOutOctets counters from the leaf switch that is connected to the server's Network Interface Controller. The ifHCInOctets shows the total number of bytes received on the interface, and the ifHCOutOctets shows the total number of bytes transmitted out of the interface. The values of these counters are measured since the initialization of the network management system.

We execute a C program, to create load to the access memory. The program allocates large blocks of memory and fills them with 1. To impose load to the CPU, we run stress tool [78] in the server, which spawns workers spinning on the sqrt(). A single worker imposes 100% load in a CPU thread. In our case, we have in total 24 physical cores with hyperthreading, hence 48 threads in total. Spawning one worker, we fully utilize one thread. Finally, iperf tool [79] is used to generate bulk connection between our server and another one which belongs in the same network. The protocol that we use is TCP.

The energy consumption of the server is collected with ipmitool [80] directly from the power supply units (PSU). The server has two PSUs, so we calculated the total power consumption by accumulating the energy output from both.

4.2. Test cases

We create 9 test cases based on different memory loads, from 0 to 100%. For each memory load, we run 8 CPU loads from 0 to 100%. The test cases are designed in such a way to cover scenarios for all the range of memory and CPU load of the server. The size of the access memory is 128 GB. Therefore, we have one test case when memory is not stressed. In this case, the only memory overhead is approximately 1.4 GB, and is produced by the system processes of the operating system. Then we have 7 more test cases, running our C program which stresses the memory up to maximum with in a pure CPU stress test and also well explain the powerconsumption. Variable R2σ% #Out

Variable	R2	σ%	#Out
host_df_df.root_free_SQ	0.992	0.216	0.289
host_df_df.root_used	0.996	0.146	0.195
perf_instructions	0.991	0.232	0.310
perf_cycles	0.991	0.233	0.310
Table IBUR N CPU	0.997	0.144	0.192

Figure 2. Burn CPU residuals for variable host_df_df.root_free_SQ. Figure 2 shows the model residuals when using a model with only the variable host_df_df.root_free_SQ as single explanatory variable. There are only a few outliers, never-theless influencing σ significantly. B. Memory Loop The dataset "Mem Loop" consists of 2017 completed data records (without N/A values). Table II shows the best variables explaining the power, Variable R2σ% #Out

Variable	R2	σ%	#Out
perf_context.switches	0.998	0.268	0.339
perf_cache.references	0.998	0.285	0.361
perf_context.switches_SQ	0.999	0.213	0.270
	18	0.196	0.248

0.977 0.919 1.163 0

Table II MEMORY LOOP. Figure 3. Loop memory residuals for variable perf_context.switches. Figure 3 shows the model residuals when using a model with only the variable perf_context.switches as single explanatory variable. Again a few outliers are visible, which however do not significantly influence the results. Also, it is obvious that the residuals do exhibit a deterministic component. In fact this deterministic dependence could be exploited by further refining the regression model, e.g., by adding a sin()-like function. It turns out that this is not really necessary since the one-variable model is already good enough such that any model extension is superfluous.

C. Network The dataset "Network" consists of 1937 complete data records (without N/A values). Table II shows the best variables explaining the power, Variable R² σ% #

Variable	R ²	σ%	#
host_interface_if_octets.eth1_tx	0.954	0.178	0.247
host_interface_if_packets.eth1_rx	0.953	0.181	0.251
host_interface_if_packets.eth1_tx	0.951	0.183	0.255
perf_cycles + 0.967	0.152	0.211	0
host_interface_if_packets.eth1_tx	0.991	0.079	0.110
perf_instructions + 0.967	0.152	0.211	0
host_interface_if_packets.eth1_tx	0.991	0.079	0.110
perf_cycles + 0.967	0.152	0.211	0
host_interface_if_octets.eth1_tx	0.991	0.077	0.108

17

Table III NET WOR K. Figure 4. Network residuals for variable host_interface_if_octets.eth1_tx. Figure 4 shows the model residuals when using a model with only the variable host_interface_if_octets.eth1_tx as single explanatory variable. This variable makes sense when keeping in mind that the workload mainly stresses the network card. Furthermore, Table III and Figure 4 show that although two variables are enough to explain the power consumption, there are a few outliers that do have a significant influence in the result. Removing only 13 out of 1937 data records increases R² to 0.98 for one, and 0.99 for two variables.

D. Tar Kernel This dataset contains only 18 complete data records. Inasmuch it was not possible to reduce the number of explanatory variables by using a full model first, and thus, searching for optimal variable combinations had to include all 330 explanatory variables (K=330). This was partially compensated by the fact that regression for a small dataset is much faster. Nevertheless, the resulting run times were much higher. Table IV and Figure 5 show the regression results. It can be seen that mainly host oriented variables perform best, e.g., number of blocked processes, disk writing load, but also the load averages.

Variable R² σ% #

Variable	R ²	σ%	#
host_processes_ps_state	0.911	0.661	0.816
blocked_value	0.911	0.660	0.816
host_disk.sda_disk_time_write_SQ	0.911	0.661	0.817
host_cpu.2_cpu.wait_value_SQ + 0.984	0.290	0.358	0.364
host_load_load_midterm_SQ	0.983	0.295	0.364
host_load_load_midterm_SQ	0.983	0.300	0.370
host_load_load_shortterm_SQ	0.370	0.370	0.370

Table IV TAR KE RNE L. Figure 5. Tar kernel residuals for variables host_cpu.2_cpu.wait_value_SQ + host_load_load_midterm_SQ.

E. Disk Read This dataset contains 2018 complete data records. As can be seen in Tables V and VI, and Figure 6 there are a few outliers that have a significant impact on the model accuracy. Removing them results in a sufficient fit though. We also computed models for the optimal triple of variables, which however did not improve the model quality significantly. We thus omitted them in Table V. Interestingly single variables do not include disk related variables, but when sequentially adding variables, pid_kB_rd_s_SQ turns up at place third, so it is definitely related to the power consumption. Table VI shows that for six variables R² reaches up to 0.98 without outliers, more variables though do not improve the accuracy any more.

F. Disk Write This dataset contains again

excellent results. Finally, Figure 10 shows the quantiles of the residuals plotted against the theoretical quantiles of normal distributions. The more straight lines the curves are, the more

Variable R²σ% #Outperf_cache.misses 0.761 0.588 0.822 0.925 0.297 0.415
 19perf_cache.misses_SQ 0.816 0.524 0.732 0.919 0.308 0.430 20host_cpu. 0.865 0.449
 0.627 00_cpu.system_value_SQ 0.932 0.296 0.414 21perf_context.switches_SQ 0.906 0.374
 0.523 00.973 0.196 0.274 13Table VIII DISK WRITE, SEQUENTIALLY ADDING SINGLE
 VARIABLES. Figure 7. Disk write residuals for variables perf_context.switches
 +perf_instructions_SQ + host_cpu.0_cpu.system_value_SQ. they follow a normal distribution. As
 can be seen, the model using the optimal triple (O3), as well as the model using 9 variables show
 a large region between -3.5 and 2.5 where they follow a normal distribution nicely. Outside
 this region, both models show differences in their tails caused by outliers. H. Global Model Applied
 to all Datasets Table XI shows how a model with 9 variable explains all data sets in total. The
 question now is how this global

Variable R²σ% #Outperf_cache.references_SQ 0.770 2.003
 2.708 0perf_cache.misses_SQ 0.767 2.017 2.728 0.773 1.987 2.686 16perf_cache.references
 0.683 2.353 3.181 0pid_usr + 0.920 1.181 1.597 0perf_cache.references_SQ 0.937 1.046 1.414
 57perf_instructions_SQ + 0.917 1.207 1.632 0perf_cache.references_SQ 0.933 1.080 1.460
 57perf_task.clock.msecs_SQ + 0.916 1.208 1.633 0perf_cache.references_SQ 0.933 1.078
 1.457 57perf_cache.references + 0.963 0.807 1.091 0perf_instructions_SQ + 0.976 0.646 0.874
 62pid_RSS_SQperf_cache.references + 0.963 0.807 1.091 0perf_instructions_SQ + 0.976
 0.646 0.874 62pid_MEM_SQperf_cache.references + 0.963 0.807 1.091 0perf_instructions_SQ
 + 0.976 0.646 0.874 62pid_VSZ_SQ Table IX ALL DATAS ETS .

Variable Coeff Coeff Err Intercept 7.056e+01 7.781e-03perf_cache.misses_SQ -8.195e-15
 1.671e-15pid_usr 8.401e-02 1.879e-04host_irq_irq.26_value 6.882e-03
 6.278e-05pid_system_SQ 3.717e-03 3.237e-05perf_context.switches 6.266e-04
 1.107e-05pid_kB_wr_s 1.737e-04 4.696e-06host_disk.sda_disk_merged_write -5.270e-04
 1.930e-05perf_cache.references 1.190e-06 1.357e-08pid_VSZ -1.601e-05 2.064e-07 Table
 X ALL DATAS ETS , REGRESSION COEFFICIENTS AND ERRORS OF THE COEFFICIENTS
 . ALL COEFFICIENTS ARE SIGNIFICANTLY DIFFERENT FROM ZERO (P-VALUE
 S<1.0E-7) Variable R²σ% #Outperf_cache.misses_SQ 0.767 2.017 2.728 0.773 1.987 2.686
 16pid_usr 0.916 1.209 1.635 0.935 1.062 1.437 56host_irq_irq.26_value 0.950 0.934 1.263
 00.971 0.706 0.955 63pid_system_SQ 0.966 0.772 1.043 0.977 0.634 0.857
 48perf_context.switches 0.976 0.651 0.881 0.986 0.491 0.663 65pid_kB_wr_s 0.980 0.593
 0.801 0.989 0.426 0.577 78host_disk. 0.983 0.541 0.732 0sda_disk_merged_write 0.989
 0.430 0.581 92perf_cache.references 0.985 0.508 0.687 0.990 0.407 0.550 77pid_VSZ 0.991
 0.402 0.544 0.995 0.282 0.382 90 Table XI ALL DATAS ETS , SEQUENTIALLY ADDING
 SINGLE VARIABLES. Figure 8. All datasets residuals for the sequential model using 9
 variables shown in Table X. model performs when applied to each individual data set. Table XII
 shows the respective regression results. As can be seen the global model performs
 outstandingly well, and explains the power consumption for every single dataset. The only data
 set having some deviances is "Disk Write", but when removing only 12 outliers, the situation
 dramatically changes, and 97% of the variance is explained, which is an excellent result. Figure 9.
 All datasets, ECDFs of the residuals, when using 1, 3, 5, 7, or 9 variables added sequentially. O3
 denotes the optimal combination of 3 variables. Figure 10. All datasets, Q-Q-plot of the residuals,
 when using 1, 3, 5, 7, or 9 variables added sequentially. O3 denotes the optimal combination of 3

variables. V. POWER CONSUMPTION OF A SINGLE PROCESS The above derived models describe the power consumption of a computer by a combination of system wide variables Y_i , $i = 1, \dots, I$, as well as variables X_{jl} , $j = 1, \dots, J$ describing individual process pl , $l = 1, \dots, L$. Let the power consumption of a computer with no load be denoted by P_0 (the intercept), and the respective coefficients of the regression model be called α_i for system wide variables, and β_j for per process variables. The linear model then relates

Data set R2


	R2	σ^2	% #Out	"Burn CPU"	0.992
0.225	0.3	0	"Mem Loop"	0.999	0.144
0.24	0	"Network"	0.944	0.196	0.27
0.0	0.974	0.134	0.19		
17	"Tar Kernel"	0.965	0.556	0.69	0
0	"Disk Read"	0.946	0.389	0.54	0.0
0.974	0.265	0.36	11	"Disk Write"	0.89
0.407	0.57	0.0	0.965	0.221	0.31
12	Table XI	APP LYING T HE G LOB AL MO DE L FRO M TABLE X I TO T HE RE SPE CT IVE DATA SET S.			

the power consumption P of a computer in the following way: $P = P_0 + \sum_{i=1}^I \alpha_i Y_i + \sum_{j=1}^J \beta_j X_{jl}$

(1) By setting $S_j = \sum_{l=1}^L X_{jl}$, (1) can be written as $P = P_0 + \sum_{i=1}^I \alpha_i Y_i + \sum_{j=1}^J \beta_j S_j$. (2) In order to derive the power caused by an individual process Pl , we assume that the relevant system wide variables are caused by the processes, which are themselves described by their individual per process variables. To verify this assumption we explained the system wide variables `host_irq_irq.26_value` and `host_disk.sda_disk_merged_write` by the per process variables. The resulting R^2 were 0.956 and 0.99, i.e., they can be almost entirely explained by per process variables. We therefore write $Y_i = \sum_{j=1}^J \gamma_{ij} S_j$, (3) and adapt (1) to a new form $P = P_0 + \sum_{i=1}^I \alpha_i \sum_{j=1}^J \gamma_{ij} S_j + \sum_{j=1}^J \beta_j S_j = P_0 + \sum_{j=1}^J (\sum_{i=1}^I \alpha_i \gamma_{ij} + \beta_j) S_j$ (4) by defining $\eta_j = \sum_{i=1}^I \alpha_i \gamma_{ij} + \beta_j$. Thus, we now describe the global power consumption by per process variables only. At this point we can derive the power P_{pl} caused by pl to be $P_{pl} = \sum_{j=1}^J \eta_j X_{jl}$. (5) It can easily be seen that $P = P_0 + \sum_{l=1}^L P_{pl}$, as it should be.

VI. CONCLUSION We defined and implemented a measurement platform at the University of Vienna which we used to measure the dependence of the power consumption of a standard PC on a synthetic workload. By using multivariate regression we explain this power consumption by a subset of 165 process and system variables we measured during our experiments. We explore this dependence in depth and demonstrate that the obtained model quality is very good. We derive a global model for all data sets that also explains each single dataset excellently. Using simple linear analysis we also define the power consumption of a single process. Although we present a multitude of results, the main result of this paper is actually the methodology itself. The described regression results merely demonstrate that our approach does yield excellent prediction results. Future works include choosing the best values to measure in order to reduce the impact of measurement on the system.

ACKNOWLEDGMENT This work was done during a short term scientific mission (Cost action 0804) of Georges Da Costa (IRIT, Toulouse, France) to the University of Vienna (Austria) from 8, January to 12, February 2010, and was



EuropeanCommission.REFERENCES[1] R. Joseph and M. Martonosi, "Run-time power estimation in high performance microprocessors," in ISLPED '01: Proceedings of the 2001 international symposium on Low power electronics and design. ACM, 2001.[2] I. Kadayif, T. Chinoda, M. Kandemir, N. Vijaykirsnan, M. J. Irwin, and A. Sivasubramaniam, "vec: virtual energy counters," in PASTE '01: Proceedings of the 2001 ACM SIGPLAN-SIGSOFT workshop on Program analysis for software tools and engineering. ACM, 2001.[3] X. Ma, M. Dong, L. Zhong, and Z. Deng, "Statistical power consumption analysis and modeling for gpu-based computing," in SOSP Workshop on Power Aware Computing and Systems (HotPower '09), 2009.[4] S. Rivoire, P. Ranganathan, and C. Kozyrakis, "A comparison of high-level full-system power models." in HotPower, F. Zhao, Ed. USENIX Association, 2008.[5] M. J. Crawley, Statistics: An Introduction using R. Wiley, 2005, ISBN 0-470-02297-3. [Online]. linear model using PMCs and processors' temperature-along with their respective square root and logarithm-for real HPC workloads. The variables are selected according to their correlation with the power-a given variable is added to the model if its inclusion on the linear model, after calibration, has a better correlation than before. ...

Effectiveness of Neural Networks for Power Modeling for Cloud and HPC: It's Worth It!

Article

May 20

Jean-Marc Pierson Leandro Fontoura Cupertino

View

Show abstract

... Based on these results, the authors proposed a model based on S_{PMC} as in Equation 12. In the same year, [34] conducted experiments to study the energy consumed by different applications. The authors used synthetic workload to stress the server and simultaneously measured the power consumption and 165 server performance indexes. ...

... However, the servers for model development, workload to stress the servers, and the power measurement technique used by these works are different. For instance, [9,20,33,34,38,61,64,75,79,80,90,104] used a single server for power model development, while [26,27,39,41,62,67,74,76,119,123] used multiple heterogeneous servers. Moreover, [34,79,80,90] used synthetic workload to stress the server, while [9, 20, 26, 27, 33, 38, 39, 41, 62, 64, 67, 74-76, 104, 119, 123] used different benchmarking applications and real-world workload traces. ...

... For instance, [9,20,33,34,38,61,64,75,79,80,90,104] used a single server for power model development, while [26,27,39,41,62,67,74,76,119,123] used multiple heterogeneous servers. Moreover, [34,79,80,90] used synthetic workload to stress the server, while [9, 20, 26, 27, 33, 38, 39, 41, 62, 64, 67, 74-76, 104, 119, 123] used different benchmarking applications and real-world workload traces. In addition, [76] used DW-6090 power meter, [90] used a power analyzer, [123] used Chroma 66202 power meter, [67] used IBM active energy manager, [61] used Voltcraft Energy Logger 4000, [26,39,75] used a AC power meter, [9] used Yokogawa WT210 power meter, [20,64] used WattsUp PRO ES power meter, [34] used a watt meter, [79] used a smart power meter, [62,119] used home-brew power meter, and [33] used a plogg power meter to measure the power consumption of the server. ...

Computing server power modeling in a data center: survey,taxonomy and performance evaluation

Preprint

Full-text available

Apr 2020

Leila Ismail

Huned Materwala

View

Show abstract

... Kansal 2 Wireless Communications and Mobile Computing CPU utilization, the number of missing last-level caches, and the number of bytes that were read and written, as discussed in [11]. Costa and Hlavacs proposed a generalized power consumption model of the server based on its component performance counters, such as CPU cycles per second, cache per second, and cache misses per second [12]. Beloglazov et al. presented a nonlinear model of servers that performs better than regression models [13]. ...

... Equation (11) is expressed to meet the reliability of the application layer user service F_i , and the number of virtual machines to be deployed on the server is l_i . Equations (12) and (13) indicate the resource constraints that need to be met to avoid server overload. Equation (14) means that to avoid server failure, resulting in unreliable user services, the number of virtual machines corresponding to the deployment of the same service by up to the user is 1 on the same server. ...

Energy Consumption Optimization of an IoT Monitoring Center Based on a Max-Min Ant Colony Algorithm

Article

Full-text available

Feb 2023 WIREL COMMUN MOB COM

Yinghua Tong

[View](#)

[Show abstract](#)

... The term N_{LLCM} represents the number of the missing LLC during T , and α_{mem} and γ_{mem} are the linear model parameters. A more generalized power consumption model is presented in [88] based on the server's components performance counters (i.e., CPU cycles per second, references to the cache per second, cache misses per second), as given in (11). Later, the power consumption model in (11) is further extended by Witkowski et al. [89] by including the CPU temperature in the model. ...

A Review of Data Centers Energy Consumption And Reliability Modeling

Article

[Full-text available](#)

Nov 2021

Kazi Main Uddin Ahmed

Math Bollen

Manuel Alvarez

[View](#)

[Show abstract](#)

... One can do it soon after procuring h_k in the Cloud data center. The value of memory access rate depends on the number of last level cache misses, and it can be obtained from the hardware performance counters at a low overhead [23,33]. Figure 1 shows the model of the Task Scheduler. ...

Energy and Makespan Aware Scheduling of Deadline Sensitive Tasks in the Cloud Environment

Article

[Full-text available](#)

Jun 2021

Anurina Tarafdar

Mukta Debnath

Sunirmal Khatua

Rajib K Das

[View](#)

[Show abstract](#)

... Da Costa et al. [DCH10] evaluate the power consumption of a PC by using performance counters, then extend the conception to predict the power consumption of single applications. Training data is collected by running several applications and synthetic benchmarks. ...

Evaluating and modeling the energy impacts of data centers, in terms of hardware / software architecture and associated environment

Thesis

Mar 2020

Yewan Wang

[View](#)

[Show abstract](#)

... Based on these results, the authors proposed a model based on S_PMC as in Equation (12). In the same year, Reference [34] conducted experiments to study the energy consumed by different applications. The authors used synthetic workload to stress the server and simultaneously measured the power consumption and 165 server performance indexes. ... Computing Server Power Modeling in a Data Center: Survey, Taxonomy, and Performance Evaluation

Article